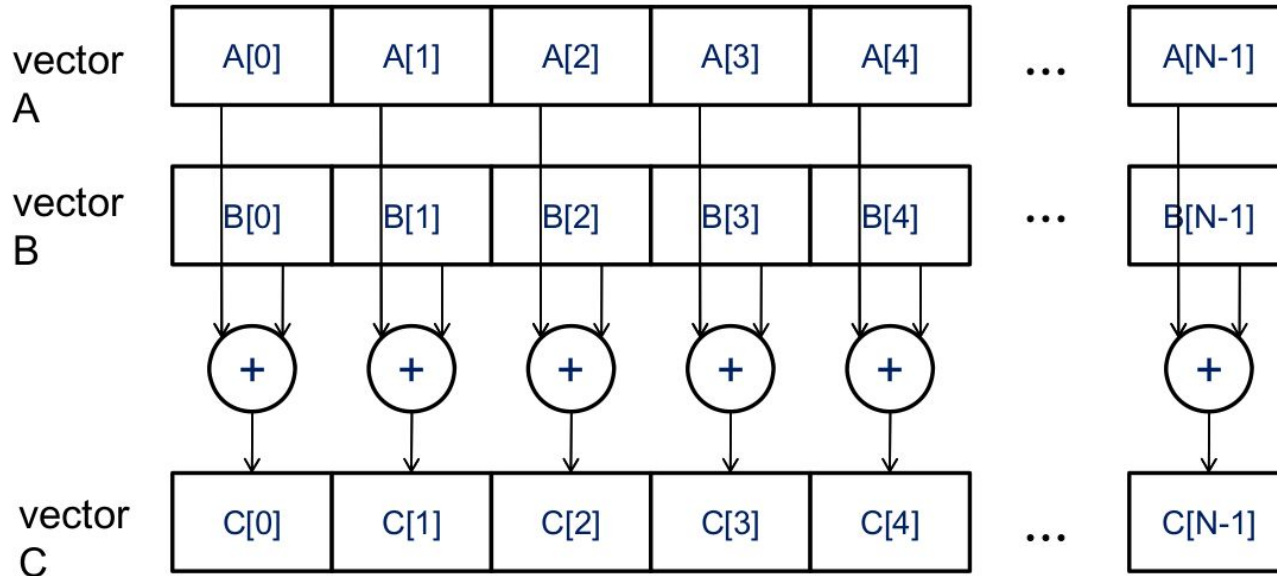


Introduction to Data Parallelism and CUDA C

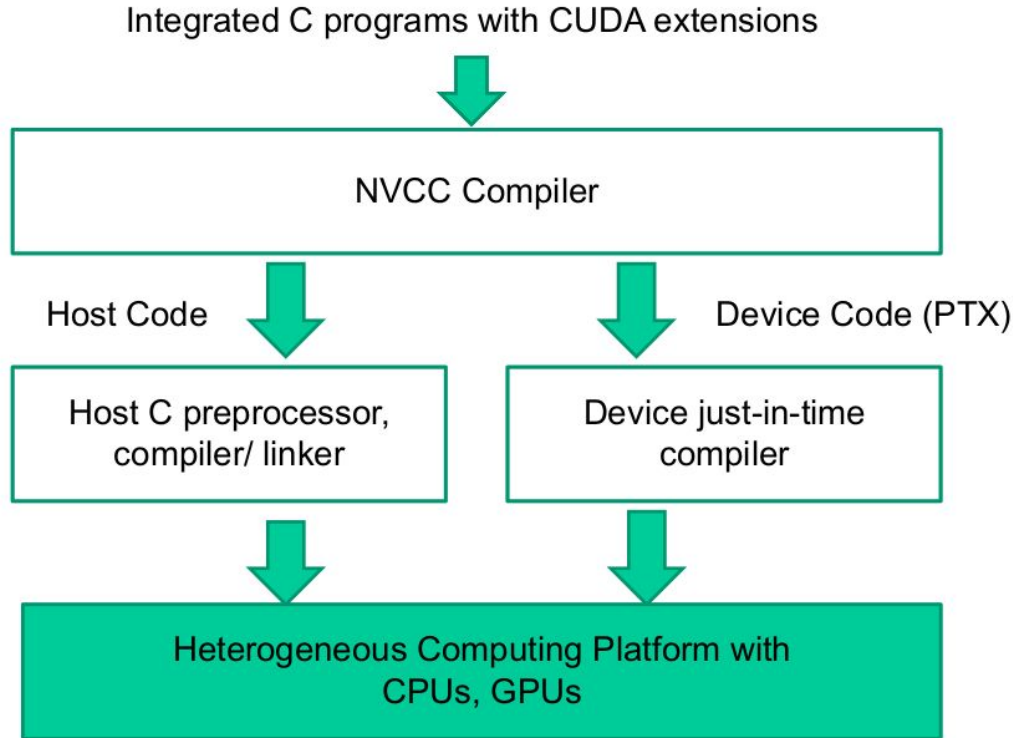
John H. Osorio Ríos



Data Parallelism



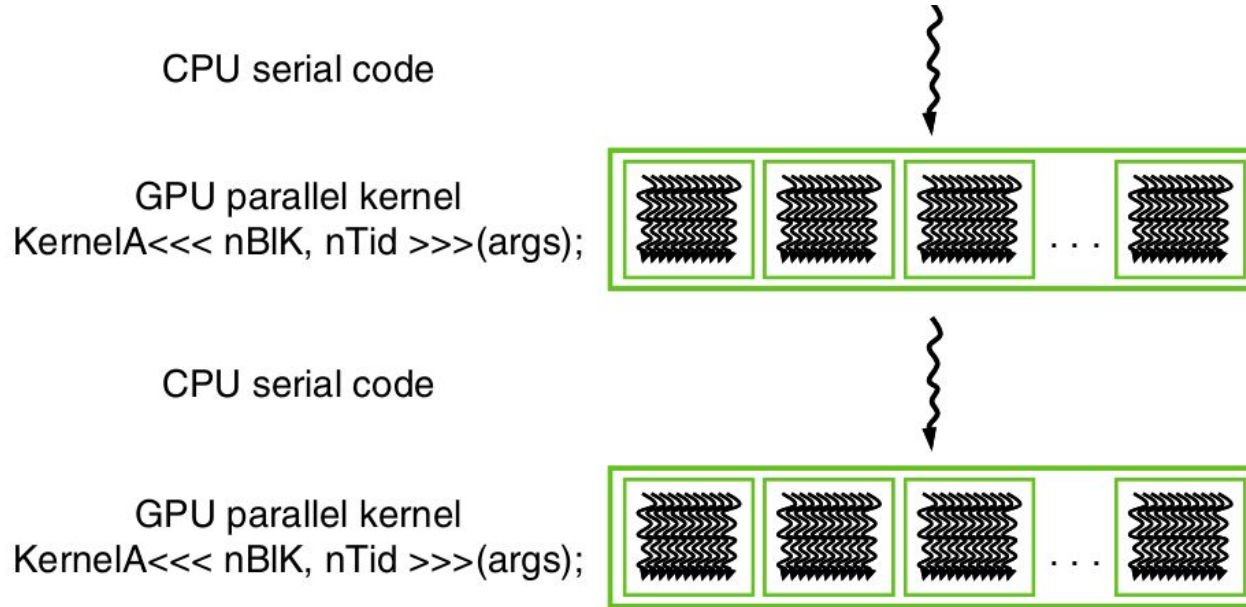
CUDA Program Structure



Compilation process



CUDA Program Structure



Execution of a CUDA program



A Vector Addition Kernel

```
// Compute vector sum h_C = h_A+h_B
void vecAdd(float* h_A, float* h_B, float* h_C, int n)
{
    for (i = 0; i < n; i++) h_C[i] = h_A[i] + h_B[i];
}

int main()
{
    // Memory allocation for h_A, h_B, and h_C
    // I/O to read h_A and h_B, N elements each
    ...
    vecAdd(h_A, h_B, h_C, N);
}
```

Traditional vector
addition



A Vector Addition Kernel

```
#include <cuda.h>
```

```
...
```

```
void vecAdd(float* A, float* B, float* C, int n)  
{
```

```
    int size = n* sizeof(float);  
    float *A_d, *B_d, *C_d;
```


```
    ...
```

```
1. // Allocate device memory for A, B, and C  
   // copy A and B to device memory
```

```
2. // Kernel launch code – to have the device  
   // to perform the actual vector addition
```

```
3. // copy C from the device memory  
   // Free device vectors
```


```
}
```



Code Scheme



Device Global Memory And Data Transfer



**Memories block
diagram**



Device Global Memory And Data Transfer

- `cudaMalloc()`
 - Allocates object in the device global memory
 - Two parameters
 - **Address of a pointer** to the allocated object
 - **Size** of allocated object in terms of bytes
- `cudaFree()`
 - Frees object from device global memory
 - **Pointer** to freed object



**CUDA API Functions to
manage device memory**



Device Global Memory And Data Transfer

```
float *d_A  
int size = n * sizeof(float);  
cudaMalloc((void**)&d_A, size);  
...  
cudaFree(d_A);
```

CUDA Malloc Example



Device Global Memory And Data Transfer

`cudaMemcpy()`

- memory data transfer
- Requires four parameters
 - Pointer to destination
 - Pointer to source
 - Number of bytes copied
 - Type/Direction of transfer



Device Global Memory And Data Transfer

```
void vecAdd(float* A, float* B, float* C, int n)
{
    int size = n * sizeof(float);
    float *d_A, *d_B, *d_C;

    cudaMalloc((void **) &d_A, size);
cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_B, size);
cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);

    cudaMalloc((void **) &d_C, size);

    // Kernel invocation code - to be shown later
    ...

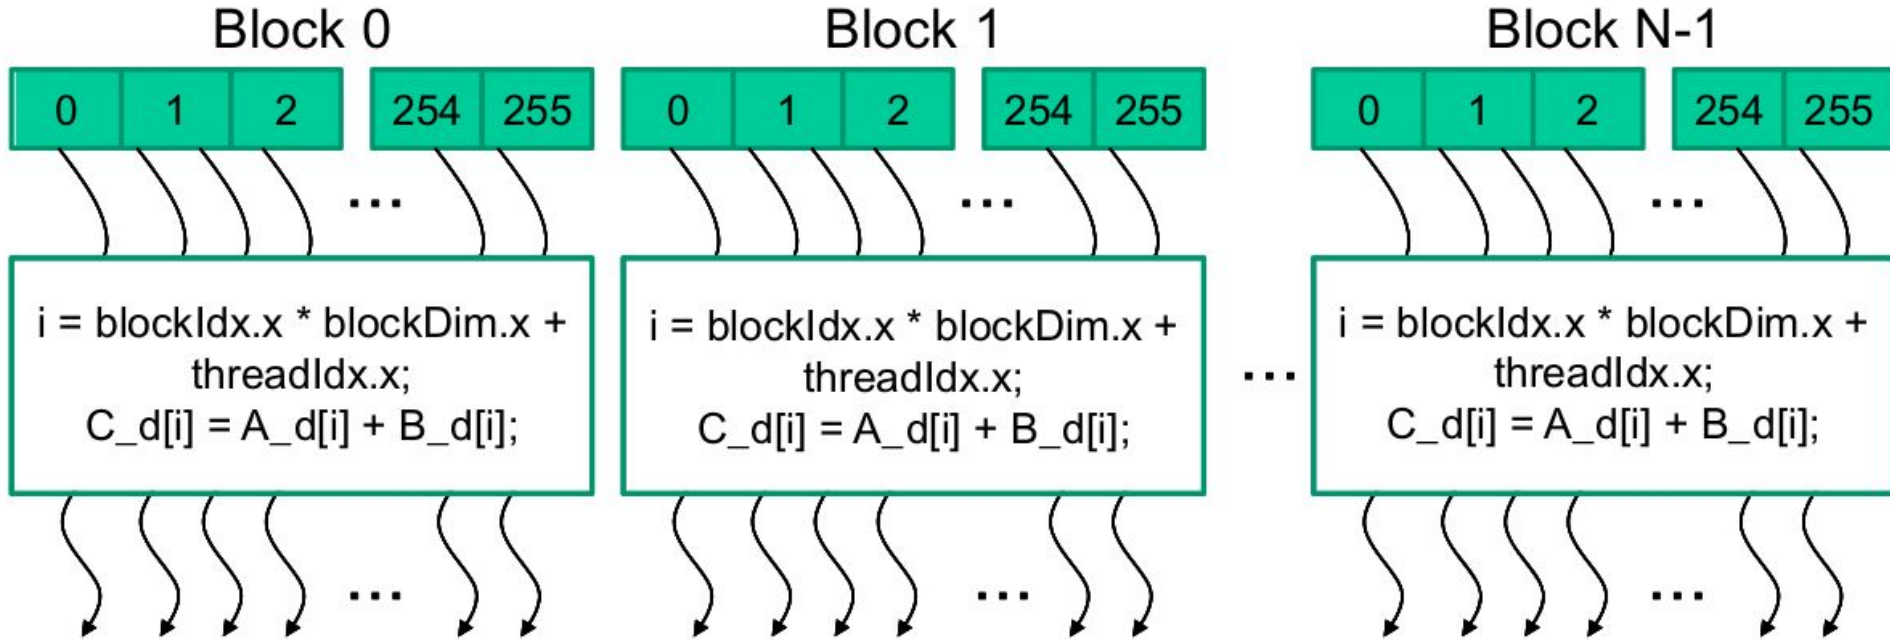
cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);

    // Free device memory for A, B, C
    cudaFree(d_A); cudaFree(d_B); cudaFree (d_C);
}
```

*A more complete
version*



Kernel Functions and Threading



Kernel Functions and Threading

```
// Compute vector sum  $C = A+B$   
// Each thread performs one pair-wise addition  
__global__  
void vecAddKernel(float* A, float* B, float* C, int n)  
{  
    int i = threadIdx.x + blockDim.x * blockIdx.x;  
    if(i<n) C[i] = A[i] + B[i];  
}
```

Vector Addition Kernel



Kernel Functions and Threading

	Executed on the:	Only callable from the:
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host
<code>__host__ float HostFunc()</code>	host	host

**CUDA C Keywords for
Function Declaration**



Kernel Functions and Threading

```
int vectAdd(float* A, float* B, float* C, int n)
{
    // d_A, d_B, d_C allocations and copies omitted
    // Run ceil(n/256) blocks of 256 threads each
    vecAddKernel<<ceil(n/256.0), 256>>>(d_A, d_B, d_C, n);
}
```

**Kernel Launch
Statement**



Kernel Functions and Threading

```
void vecAdd(float* A, float* B, float* C, int n)
{
    int size = n * sizeof(float);
    float *d_A, *d_B, *d_C;

    cudaMalloc((void **) &d_A, size);
    cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_B, size);
    cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);

    cudaMalloc((void **) &d_C, size);

    vecAddKernel<<ceil(n/2560), 256>>>(d_A, d_B, d_C, n);

    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);
    // Free device memory for A, B, C
    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}
```

Complete VecAdd





THANKS

john@sirius.utp.edu.co

