


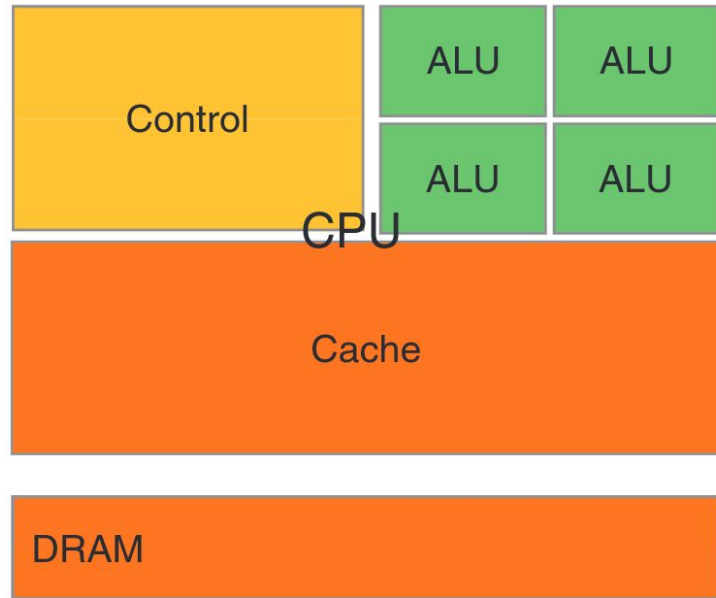
# Vector Addition

## Simple Example

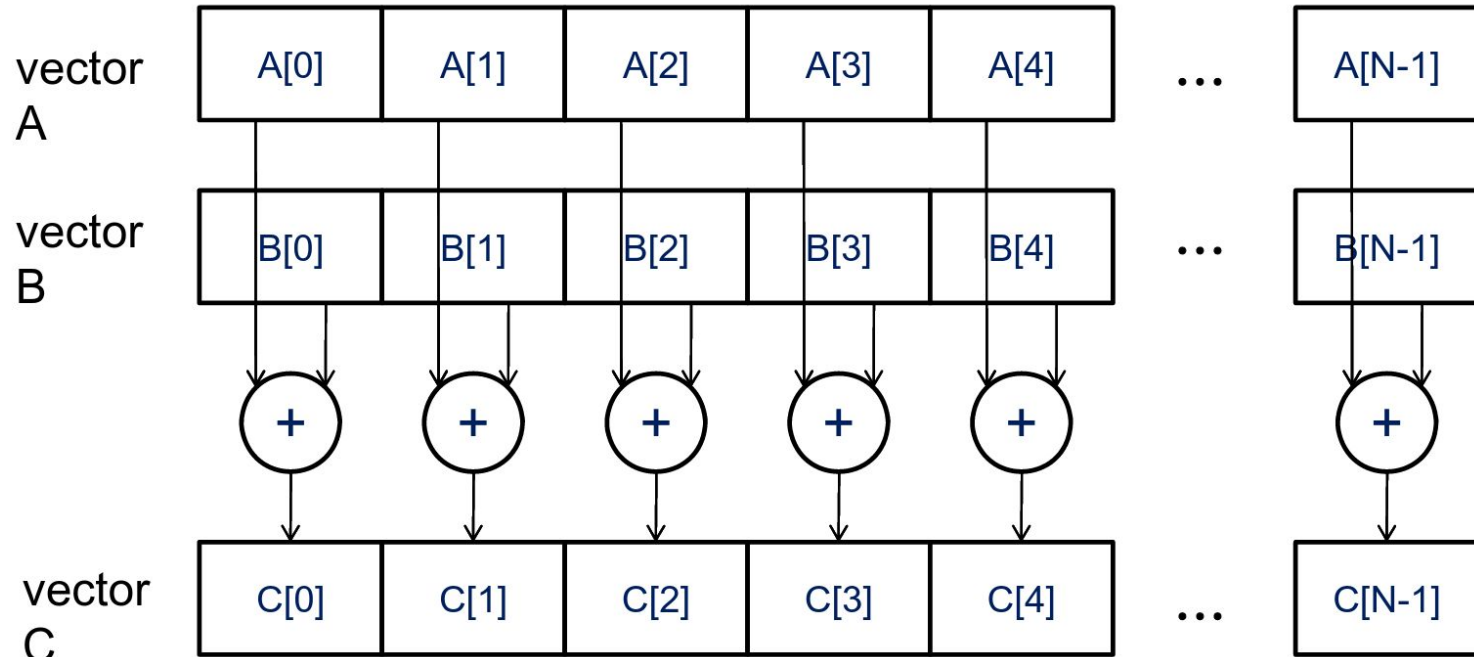
John H. Osorio Ríos



# Overview (1/3)




# Overview (2/3)



# Overview (3/3)


CPU serial code

GPU parallel kernel  
`KernelA<<< nBlk, nTid >>>(args);`



CPU serial code

GPU parallel kernel  
`KernelA<<< nBlk, nTid >>>(args);`



# Traditional C Vector Addition

```
// Compute vector sum h_C = h_A+h_B
void vecAdd(float* h_A, float* h_B, float* h_C, int n)
{
    for (i = 0; i < n; i++) h_C[i] = h_A[i] + h_B[i];
}

int main()
{
    // Memory allocation for h_A, h_B, and h_C
    // I/O to read h_A and h_B, N elements each
    ...
    vecAdd(h_A, h_B, h_C, N);
}
```




# vecAdd Process


```
#include <cuda.h>
...
void vecAdd(float* A, float*B, float* C, int n)
{
    int size = n* sizeof(float);
    float *A_d, *B_d, *C_d;
    ...
    1. // Allocate device memory for A, B, and C
       // copy A and B to device memory

    2. // Kernel launch code – to have the device
       // to perform the actual vector addition

    3. // copy C from the device memory
       // Free device vectors
}
```



# Host and Device Memory



# cudaMalloc and cudaFree (1/2)

- cudaMalloc()
  - Allocates object in the device global memory
  - Two parameters
    - **Address of a pointer** to the allocated object
    - **Size** of allocated object in terms of bytes
- cudaFree()
  - Frees object from device global memory
    - **Pointer** to freed object





## cudaMalloc and cudaFree (2/2)

```
float *d_A  
int size = n * sizeof(float);  
cudaMalloc((void**)&d_A, size);  
...  
cudaFree(d_A);
```



# cudaMemcpy

cudaMemcpy()

- memory data transfer
- Requires four parameters
  - Pointer to destination
  - Pointer to source
  - Number of bytes copied
  - Type/Direction of transfer



# More Complete Version

```
void vecAdd(float* A, float* B, float* C, int n)
{
    int size = n * sizeof(float);
    float *d_A, *d_B, *d_C;

    cudaMalloc((void **) &d_A, size);
    cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_B, size);
    cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);

    cudaMalloc((void **) &d_C, size);

    // Kernel invocation code - to be shown later
    ...

    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);

    // Free device memory for A, B, C
    cudaFree(d_A); cudaFree(d_B); cudaFree (d_C);
}
```

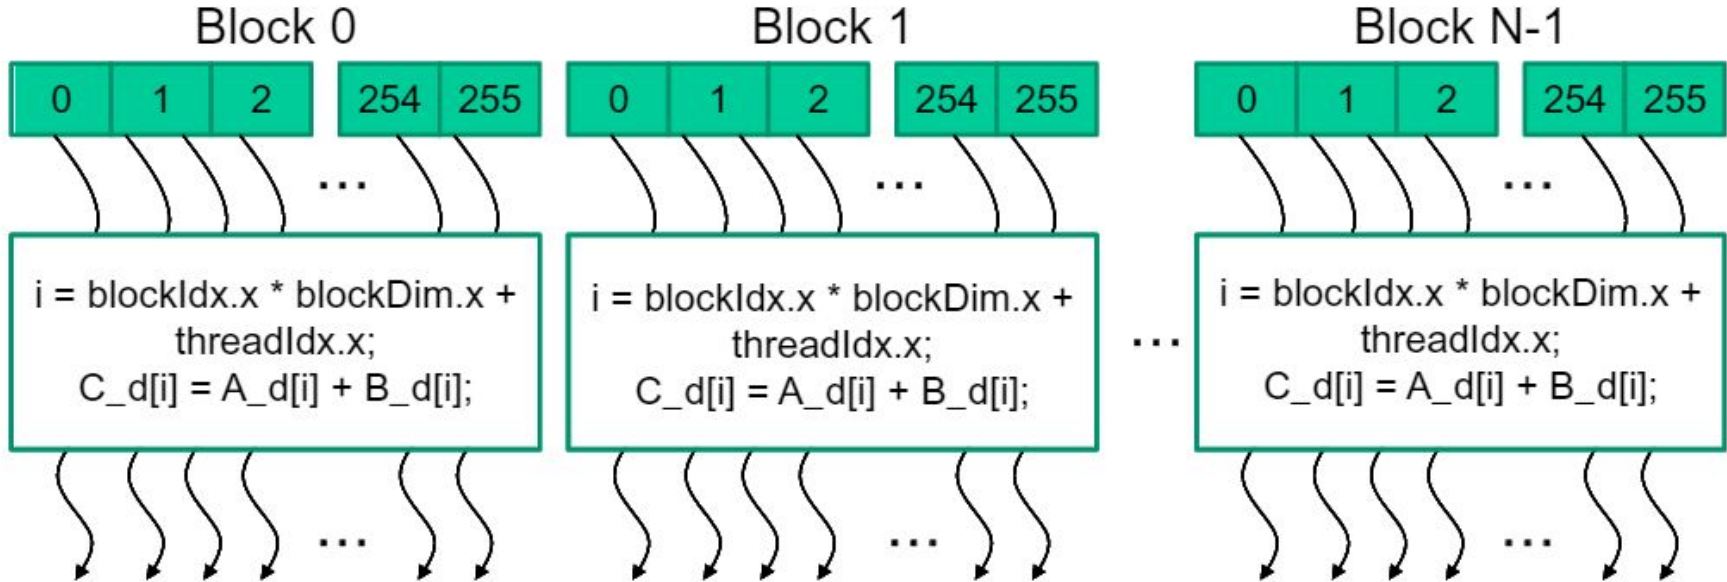


# Error Handling

```
cudaError_t err = cudaMalloc((void **) &d_A, size);  
  
if (err != cudaSuccess) {  
    printf("%s in %s at line %d\n", cudaGetErrorString( err),  
        __FILE__, __LINE__);  
    exit(EXIT_FAILURE);  
}
```



# Threads Organization



# Vector Addition Kernel

```
// Compute vector sum  $C = A+B$   
// Each thread performs one pair-wise addition  
__global__  
void vecAddKernel(float* A, float* B, float* C, int n)  
{  
    int i = threadIdx.x + blockDim.x * blockIdx.x;  
    if(i<n) C[i] = A[i] + B[i];  
}
```



# Kernel Launch

```
int vectAdd(float* A, float* B, float* C, int n)
{
    // d_A, d_B, d_C allocations and copies omitted
    // Run ceil(n/256) blocks of 256 threads each
    vecAddKernel<<<ceil(n/256.0), 256>>>(d_A, d_B, d_C, n);
}
```



# CUDA C Keywords Function Declaration

	Executed on the:	Only callable from the:
<code>__device__</code> float DeviceFunc()	device	device
<code>__global__</code> void KernelFunc()	device	host
<code>__host__</code> float HostFunc()	host	host





# Complete Function

```
void vecAdd(float* A, float* B, float* C, int n)
{
    int size = n * sizeof(float);
    float *d_A, *d_B, *d_C;

    cudaMalloc((void **) &d_A, size);
    cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_B, size);
    cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);

    cudaMalloc((void **) &d_C, size);
    vecAddKernel<<<ceil(n/2560), 256>>>(d_A, d_B, d_C, n);

    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);
    // Free device memory for A, B, C
    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}
```



# TODO (1/1)

- Make your own code to add vectors
- Try to measure times
- Use error handling - “Saves a lot of work”
- ENJOY :)



# Bibliography (1/1)

- Programming Massively Parallel Processors



# THANKS

[john@sirius.utp.edu.co](mailto:john@sirius.utp.edu.co)

